



Full YOLOv8n inference comparison on Jetson Orin Nano GPU

FP32 FP16 INT8 PyTorch FP32 baseline

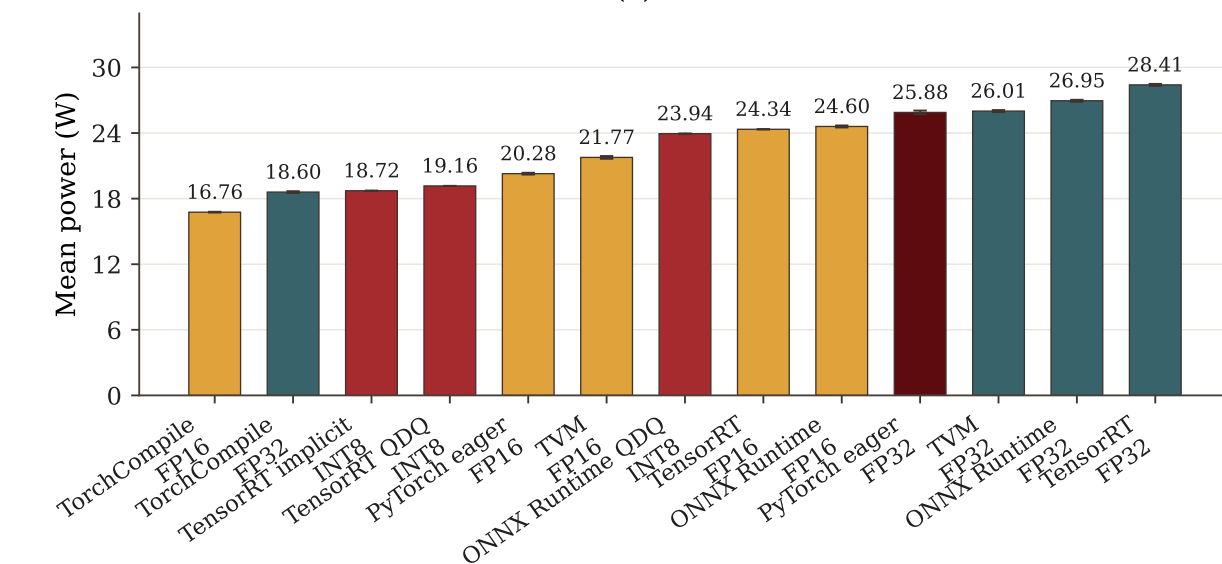
(a) Latency



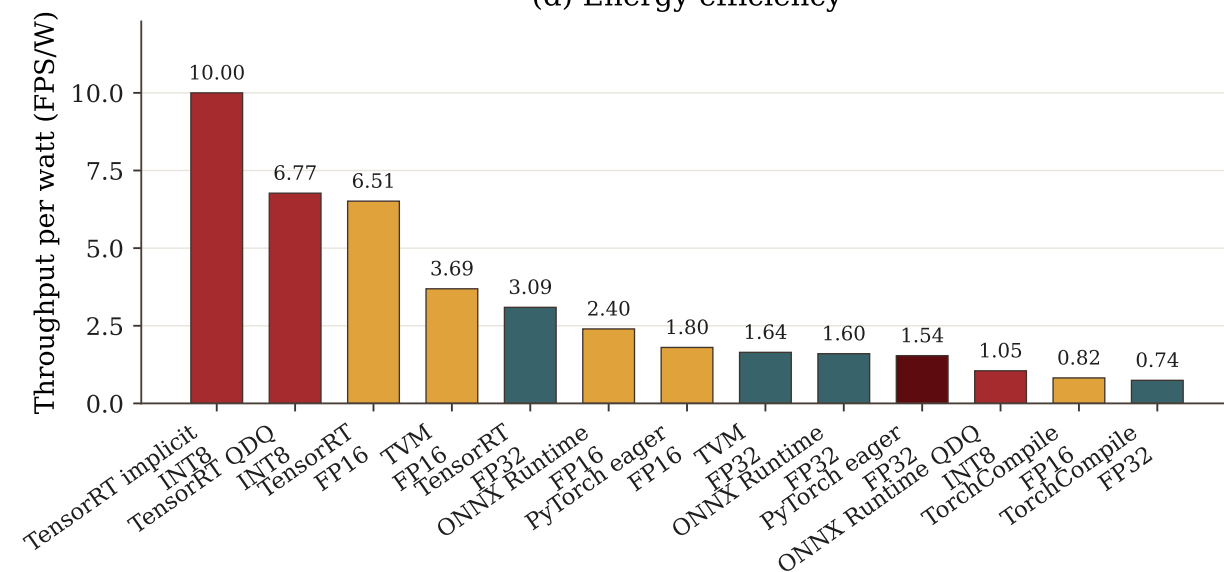
(b) Normalized speedup



(c) Power



(d) Energy efficiency



Only successful runs are included. Speedup is normalized to PyTorch eager FP32 on the same platform. Latency error bars show $\pm$ half-width of the 95% confidence interval where available.