



FP32 vs FP16 YOLOv8n inference on Jetson Orin Nano GPU

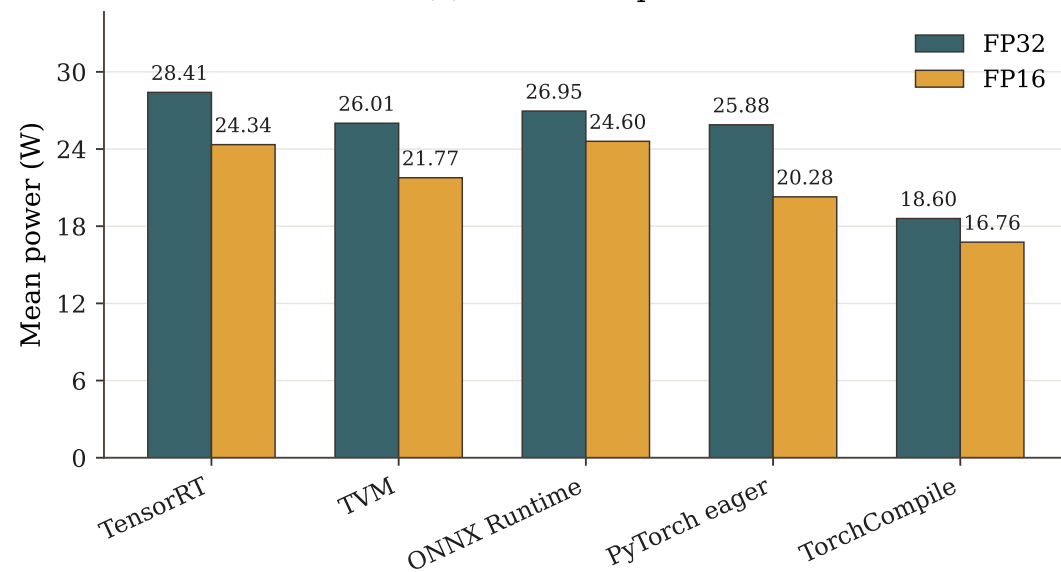
(a) Latency comparison



(b) FP16 latency speedup



(c) Power comparison



(d) Energy efficiency comparison



Only backends with successful FP16 runs are included. Latency error bars show $\pm$ half-width of the 95% confidence interval.