



INT8 YOLOv8n inference on Jetson Orin Nano GPU


(a) INT8 latency




(b) Speedup over FP32 baseline



(c) Power



(d) Energy efficiency



Speedup is computed relative to the PyTorch FP32 eager baseline on the same platform.